

Experiment 2 Report: Evaluation Function Comparison

1. Objective

Experiment 2 compares three evaluation functions under the same search framework:

Alpha-Beta pruning + move ordering

This experiment is not a Minimax vs Alpha-Beta comparison. Experiment 1 already verified that Alpha-Beta pruning is effective. Here, the search framework is fixed, and only the evaluation function changes.

This is also a fresh rerun because Eval A was modified. The old Experiment 2 folder and generated results were discarded before rebuilding this experiment, so the CSV files, plots, and interpretation here replace the previous Experiment 2 data.

The goal is to compare how Eval A, Eval B, and Eval C affect:

- best move selection;
- best score;
- explored nodes;
- Alpha-Beta cutoffs;
- runtime;
- candidate count;
- decision differences between evaluation functions.

2. Compared Evaluations

The experiment uses the current implementations in `ai/evaluation.py`.

Evaluation	Current function	Main idea
Eval A	<code>eval_basic</code>	Updated basic defensive evaluation based on consecutive segments plus extra defensive open-segment scoring.
Eval B	<code>eval_intermediate</code>	Open / blocked shape evaluation based on segment length and open ends.
Eval C	<code>eval_advanced</code>	Five-cell window potential evaluation with a small center bonus.

One important detail is that the current Eval A is not a pure minimal consecutive-length evaluator anymore. It now calls both consecutive-segment logic and defensive open-segment logic. Therefore, Eval A can be slower than Eval B even though it is conceptually the basic evaluation.

3. Controlled Conditions

To make the comparison fair, all evaluations use the same:

- board size: 15 x 15 ;
- search algorithm: Alpha-Beta + move ordering;
- candidate generator: `generate_candidate_moves(board, radius=2)` ;
- search depths: 1, 2, 3 ;
- fixed test positions;
- player to move for each position;
- move ordering implementation.

No depth 4 is tested. No random positions are used.

Each search result records:

- best move;
- best score;
- explored nodes;
- cutoffs;
- time in milliseconds;
- candidate count;
- whether the move is legal;
- whether the board stayed unchanged after search.

Detailed results are saved in:

- `results/evaluation_comparison.csv`

Aggregated results are saved in:

- `results/evaluation_summary.csv`
- `results/evaluation_comparison_summary.csv`
- `results/evaluation_comparison_summary.md`

Best-move comparison is saved in:

- `results/evaluation_best_move_diff.csv`

4. Fixed Positions

The fixed positions are defined in `positions.py` and exported as readable board states in:

- `results/evaluation_positions.txt`

The experiment uses six fixed positions:

Position	Purpose
E1_length_pressure	Tests immediate length pressure against more open development space.
E2_open_vs_blocked	Contrasts blocked longer-looking lines with open shapes.
E3_window_potential	Gives Eval C a chance to prefer multi-window potential.
E4_attack_defense	Compares attack choices with defensive potential reduction.
E5_center_space	Compares edge-side local extension with central development.
E6_complex_midgame	Uses a compact midgame with several two- and three-stone patterns.

The average candidate count is `57.0` , and the same candidate generator is used for all three evaluations.

5. Correctness Checks

The generated raw CSV contains 54 rows:

$6 \text{ positions} \times 3 \text{ depths} \times 3 \text{ evaluations} = 54 \text{ runs}$

The rerun checks passed:


- illegal moves: 0 ;
- polluted board states after search: 0 ;
- tested depths: 1, 2, 3 ;
- search method: Alpha-Beta + ordering ;
- evaluations: Eval A , Eval B , Eval C .

This confirms that all three evaluations are called through the same search interface and use the same candidate move generator.

6. Summary Table

Depth	Evaluation	Positions	Avg time ms	Avg nodes	Avg cutoffs	Avg candidate count
1	Eval A	6	35.2078	58.0	0.0	57.0
1	Eval B	6	24.8029	58.0	0.0	57.0
1	Eval C	6	245.8648	58.0	0.0	57.0
2	Eval A	6	1353.6469	190.5	55.67	57.0
2	Eval B	6	943.3087	211.33	55.33	57.0
2	Eval C	6	8657.382	209.5	55.33	57.0
3	Eval A	6	6419.806	4196.17	145.5	57.0
3	Eval B	6	4527.5326	4154.17	150.5	57.0
3	Eval C	6	37228.8886	4151.83	136.67	57.0

7. Runtime Comparison



Runtime increases strongly as search depth increases. Eval C is consistently the most expensive evaluation because it scans five-cell windows and includes more board-wide potential information.

At depth 3:

Eval A average time: 6419.8060 ms

Eval B average time: 4527.5326 ms

Eval C average time: 37228.8886 ms

Eval C is about 8.22x slower than Eval B at depth 3:

$$37228.8886 / 4527.5326 \approx 8.22$$

The updated Eval A is also slower than Eval B in this rerun. This is not a measurement error. In the current implementation, Eval A does more work than Eval B:

Eval A:

```
collect_segments(player)
collect_segments(opponent)
collect_open_segments(player)
collect_open_segments(opponent)
```


Eval B:

```
collect_open_segments(player)
collect_open_segments(opponent)
```


So although Eval A is conceptually the basic evaluation, the updated defensive version performs more scans than Eval B.

The log-scale runtime chart is also generated for readability:

Average time (ms) vs depth



8. Node Comparison



All three evaluations use the same candidate generator and the same Alpha-Beta implementation. Therefore, the node counts are much closer than the runtime numbers.

At depth 1, node counts are identical:

Eval A average nodes: 58.0

Eval B average nodes: 58.0

Eval C average nodes: 58.0

This is expected because depth 1 evaluates the same candidate set without meaningful pruning.

At depth 3:


Eval A average nodes: 4196.17

Eval B average nodes: 4154.17


Eval C average nodes: 4151.83

The difference is small. This means the main runtime gap is caused by the cost of each evaluation call, not by Eval C exploring a much larger tree.

The log-scale node chart is also generated:



9. Cutoffs



Different evaluations can still affect Alpha-Beta pruning indirectly. The evaluation score influences move ordering and alpha/beta boundary updates.

At depth 3:

Eval A average cutoffs: 145.50

Eval B average cutoffs: 150.50

Eval C average cutoffs: 136.67

Eval B has the highest average cutoff count at depth 3 in this run. Eval C has fewer cutoffs, but it also has roughly the same node count as Eval B. Therefore, the lower Eval C speed is mainly explained by expensive evaluation logic rather than by worse pruning alone.

10. Best Move Differences

Best-move differences are recorded in:

- `results/evaluation_best_move_diff.csv`

Across 18 position-depth cases:

same best move: 11 / 18

different best move: 7 / 18

This is useful for the report because Experiment 2 is not only about runtime. It also shows that different evaluation ideas can lead to different decisions under the same search framework.

Representative cases:

Position	Depth	Eval A move	Eval B move	Eval C move	Interpretation
E1_length_pressure	2	7 9	8 8	8 8	Eval A keeps the local extension, while Eval B/C prefer another shape.
E1_length_pressure	3	9 6	9 9	9 9	Deeper search still separates Eval A from Eval B/C.
E3_window_potential	2	5 9	5 9	9 5	Eval C selects a different potential-window move.
E5_center_space	1	4 6	4 6	6 6	Eval C favors central / spatial potential earlier.
E5_center_space	2	4 6	6 6	6 6	Eval B/C agree on the development move while Eval A extends locally.
E5_center_space	3	6 6	9 9	4 6	All three evaluations disagree in the center-space position.

Position	Depth	Eval A move	Eval B move	Eval C move	Interpretation
E6_complex_midgame	2	6 6	8 9	8 9	Eval B/C agree in the compact midgame while Eval A differs.

Several depth-3 rows return 100000.0 scores because the search finds a terminal winning line under that evaluation horizon. In those cases, identical scores and moves are expected even if the static evaluation functions differ.

11. Interpretation

The rerun supports four main observations.

First, Eval A changed meaning after the recent update. It still represents the basic evaluation family in this project, but the current implementation includes extra defensive open-segment scoring. This makes it more defensive and more expensive than a pure consecutive-length evaluator.

Second, Eval B is the best runtime trade-off in this dataset. It captures open/blocked shape information, and it is faster than the updated Eval A because it directly uses one open-segment scan per side.

Third, Eval C produces some different decisions, especially in positions designed around center space and five-cell window potential. However, it is much more expensive. At depth 3, Eval C is over eight times slower than Eval B while exploring nearly the same number of nodes.

Fourth, more complex evaluation does not automatically mean better practical AI at shallow depth. Eval C may encode richer long-term information, but under depths 1 to 3 its runtime cost is very high, and the decision differences should be weighed against that cost.

12. Conclusion

Experiment 2 shows that evaluation choice affects both decision behavior and runtime, even when the search algorithm is fixed.

- Eval A is now a defensive basic evaluator, but not the fastest implementation.
- Eval B provides the strongest practical balance between shape awareness and runtime.
- Eval C captures five-cell window potential but is much slower.
- Best moves differ in 7 / 18 cases, so the evaluations are not merely different weight settings.
- Node counts are similar across evaluations, so runtime differences mainly come from evaluation cost.

For later AI configuration selection, these results suggest that Eval B is a strong candidate for practical play, while Eval C should only be used if its decision quality clearly justifies the extra time.

13. How to Reproduce

Run the full comparison:

```
python experiments/experiment2_evaluations/compare_evaluations.py
```

Export readable positions:

```
python experiments/experiment2_evaluations/export_positions.py
```

Generate plots:

```
python experiments/experiment2_evaluations/plot_evaluation_comparison.py
```

If plotting dependencies are missing:

```
python -m pip install -r requirements.txt
```